

Non-Singularity of the Gradient Descent Map for Neural Networks with Piecewise Analytic Activations

Alexandru Crăciun Debarghya Ghoshdastidar

Presented by Ian Chi, slides made with assistance by Gemini 2.1

39th Conference on Neural Information Processing Systems (NeurIPS 2025)

May 26, 2026

Outline

- 1 Introduction & Optimization Dynamics
- 2 Limitations of Classical Theory
- 3 Mathematical Preliminaries
- 4 Formal Results: Non-Singularity of the GD Map
- 5 Proof Outline: The Analytic Case
- 6 Proof Outline: The Piecewise Analytic Case
- 7 Extensions to Modern Architectures
- 8 Conclusion

The Central Problem: Training Deep Networks

- Let $f_\theta : \mathbb{R}^d \rightarrow \mathbb{R}^k$ be a neural network parameterized by weights and biases $\theta \in \mathbb{R}^p$.
- For a dataset $\{(x_i, y_i)\}_{i=1}^N$ and a loss function ℓ , the empirical risk is:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \ell(f_\theta(x_i), y_i)$$

- The Gradient Descent (GD) update rule is defined as a discrete-time dynamical system:

$$\theta_{t+1} = G_\eta(\theta_t) := \theta_t - \eta \nabla L(\theta_t)$$

where $\eta > 0$ is the step-size (learning rate) and $G_\eta : \mathbb{R}^p \rightarrow \mathbb{R}^p$ is the **GD map**.

Saddle Points in High Dimensions

- The loss landscape of overparameterized neural networks is highly non-convex.
- Critical points ($\nabla L(\theta) = 0$) can be local minima, local maxima, or saddle points.
- In high-dimensional spaces ($p \gg 1$), saddle points are exponentially more prevalent than local minima (Dauphin et al., 2014).
- **Key Question:** Why doesn't GD get trapped in these ubiquitous saddle points?

The Classical Assumption: Lipschitz Smoothness

- Previous convergence guarantees fundamentally rely on $L(\theta)$ being globally Lipschitz smooth.
- That is, $\|\nabla L(\theta) - \nabla L(\theta')\| \leq H\|\theta - \theta'\|$ for all θ, θ' .
- **Failure: Absence of Global Lipschitz Smoothness.**
- Realistic neural networks, particularly deep networks with Cross-Entropy loss, are *not* globally Lipschitz smooth. The gradients can become arbitrarily steep.

Measure Zero and Non-Singularity

Definition (Measure Zero)

A set $A \subset \mathbb{R}^p$ has Lebesgue measure zero if, for every $\epsilon > 0$, A can be covered by a countable sequence of p -dimensional hyperrectangles whose total volume is less than ϵ .

Definition (Singular vs. Non-Singular Map)

A measurable map $G : \mathbb{R}^p \rightarrow \mathbb{R}^p$ is **non-singular** if, for any set $A \subset \mathbb{R}^p$ of Lebesgue measure zero, its preimage $G^{-1}(A) = \{x \in \mathbb{R}^p \mid G(x) \in A\}$ also has measure zero.

- **Intuition:** A singular map can collapse a set of positive measure (like an entire continuous interval) into a single point or a lower-dimensional set (which has measure zero).
- If G_η is singular, a huge region of the parameter space could be mapped directly onto a saddle point in one step.

Analytic and Piecewise Analytic Functions

Definition (Analytic Function)

A function $f : U \rightarrow \mathbb{R}$ defined on an open set $U \subset \mathbb{R}^p$ is real analytic if, for every $x_0 \in U$, $f(x)$ can be represented by a convergent power series in a neighborhood of x_0 .

- **Property of Zeros:** If f is analytic and not identically zero on a connected domain, the set of its roots $\{x \mid f(x) = 0\}$ has Lebesgue measure zero.

Definition (Piecewise Analytic Function)

A function is piecewise analytic if the domain can be partitioned into a countable collection of regions (strata), where the boundaries have measure zero, and the function is strictly analytic within the interior of each region.

- Includes ReLU, Leaky ReLU, Step functions, etc.
- Note: this condition of analyticity is strictly stronger than C^∞ (consider $\exp\left(\frac{1}{-x^2}\right)$)

The Main Theorem (Crăciun & Ghoshdastidar, 2025)

Theorem (Non-Singularity for Realistic Networks)

Consider a deep neural network consisting of fully connected, convolutional, or attention layers. Let the non-linear activations in the layers be piecewise analytic (e.g., ReLU, Sigmoid, Tanh, Leaky ReLU). Fix any dataset and any analytic loss function ℓ (e.g., MSE, Cross-Entropy).

Then, for **almost all** step-sizes $\eta > 0$, any (Stochastic) Gradient Descent map G_η is **non-singular**.

- **Significance:** Removes the Lipschitz assumption completely.
- Holds for practical architectures.
- Holds for arbitrarily large (but finite) step-sizes, except for a set of measure zero in the space of η .
- note: non-singularity is a different condition than what we normally wish for, that being convergence.

Connecting Non-Singularity to Optimization Theory

- With Theorem 1 established, we can strictly apply center-stable manifold theorems to deep ReLU networks.
- **Saddle Point Avoidance:** If the set of strict saddle points has measure zero, the set of initial weights θ_0 that eventually converge to a saddle point also has measure zero.
- **Stochastic Gradient Descent (SGD):** The theorem explicitly covers the SGD map as well. If the loss is evaluated on a mini-batch, the resulting map is still non-singular.

Step 1: The Jacobian and Diffeomorphisms

- Assume for a moment the activation functions are purely analytic (e.g., Sigmoid, Tanh, Softplus).
- The loss landscape $L(\theta)$ is an analytic function of θ .
- The GD map is $G_\eta(\theta) = \theta - \eta \nabla L(\theta)$.
- The Jacobian matrix is $JG_\eta(\theta) = I - \eta \nabla^2 L(\theta)$.
- By the **Inverse Function Theorem**, if $\det(JG_\eta(\theta)) \neq 0$ at a point θ , then G_η is a *local diffeomorphism* around θ .

Step 2: Analyzing the Zeros of the Determinant

Lemma (Local Diffeomorphisms Preserve Measure Zero)

Let $S = \{\theta \in \mathbb{R}^p \mid \det(JG_\eta(\theta)) = 0\}$ be the set of critical points of the map G_η . If S has Lebesgue measure zero, then G_η is non-singular.

- *Proof Intuition:* Outside of S , the space can be covered by countably many open sets where G_η is a diffeomorphism. Since diffeomorphisms map measure zero sets to measure zero sets, any measure zero set pre-imaged outside S remains measure zero. Since S itself has measure zero, the whole preimage has measure zero.
- **Goal:** Prove that $S = \{\theta \mid \det(I - \eta \nabla^2 L(\theta)) = 0\}$ has measure zero.

Step 3: Eigenvalues of the Hessian

- The condition $\det(I - \eta \nabla^2 L(\theta)) = 0$ implies that $\nabla^2 L(\theta)$ has an eigenvalue exactly equal to $\frac{1}{\eta}$.
- Let $\lambda_i(\theta)$ be the eigenvalues of the Hessian $\nabla^2 L(\theta)$.
- **Case A (Trivial):** All eigenvalues $\lambda_i(\theta)$ are strictly constant everywhere in the parameter space. In this case, for any η not exactly equal to $1/\lambda_i$, the determinant is never zero. The map is globally non-singular.
- **Case B (Realistic):** The eigenvalues $\lambda_i(\theta)$ change depending on θ .

Step 4: Analytic Paths and Eigenvalue Parameterization

- Take any two points θ_1, θ_2 . Because $\mathbb{R}^p$ is connected, we can interpolate them with an analytic path $\gamma(t)$ such that $\gamma(0) = \theta_1$ and $\gamma(1) = \theta_2$.
- The composition $L(\gamma(t))$ is analytic. The Hessian evaluated along this path, $H(t) = \nabla^2 L(\gamma(t))$, is an analytic matrix-valued function.
- **Kato's Theorem (1995):** The eigenvalues of an analytically parameterized symmetric matrix can be globally parameterized by analytic functions of t .
- Thus, $\lambda_i(t)$ are analytic functions.

Step 5: Concluding the Analytic Case

- Because $\lambda_i(t)$ is an analytic function, if it is non-constant, the set of times t where $\lambda_i(t) = \frac{1}{\eta}$ is discrete (and therefore has measure zero).
- Extending this from paths to the full parameter space $\mathbb{R}^p$: an analytic function that is not strictly constant can only equal a specific value $\frac{1}{\eta}$ on a set of measure zero.
- Therefore, $S = \{\theta \mid \det(JG_\eta(\theta)) = 0\}$ has measure zero for any chosen η .
- By Lemma 1, the GD map G_η is non-singular.

The Challenge of Piecewise Activations

- For activations like ReLU, $L(\theta)$ is **not** analytic. It is not even globally differentiable.
- The Hessian $\nabla^2 L(\theta)$ is undefined at the "kinks" (e.g., when a pre-activation is exactly 0).
- The composition of "almost everywhere analytic" functions is *not* necessarily "almost everywhere analytic".
- Applying the chain rule blindly across layers fails.


Stratification of the Parameter Space

- The authors circumvent this by studying the **loss landscape geometry directly**.
- For fixed data X , the piecewise nature of the network partitions the parameter space $\mathbb{R}^p$ into a set of countably disjoint regions: $\Omega_1, \Omega_2, \dots, \Omega_K$.
- **Interior Regions:** Inside each open region Ω_j , no neuron's pre-activation crosses zero. The network behaves exactly like a purely analytic network.
- **Boundaries:** The boundaries $\partial\Omega_j$ are the manifolds where at least one pre-activation is exactly zero.

Handling the Boundaries

- Inside each Ω_j , the restricted map $G_\eta|_{\Omega_j}$ is analytic. By the previous proof, it is non-singular on the interior.
- The critical step is proving that the union of all boundaries (the problematic points), $B = \bigcup \partial\Omega_j$, has measure zero, **and** that the image $G_\eta(B)$ does not somehow expand to have positive measure.
- It is inductively proven that the pre-activations are analytic functions of the parameters up to the kink, the sets where they equal zero form lower-dimensional analytic manifolds.

Handling the Boundaries (cont.)



In **red** are the points where the network map is **not analytic** ($S(f_{\theta})$).

The complement of the **green** set are the points where the proof guarantees that the network map **is analytic** ($\mathbb{R}^2 \setminus B(f_{\theta})$). We always have $S(f_{\theta}) \subseteq B(f_{\theta})$.

Figure: Illustration taken from paper

Beyond Fully Connected Networks

- **Convolutional Neural Networks (CNNs):**
- Convolutions are strictly linear operations. The composition of linear operations with piecewise analytic activations follows the exact same stratification rules as dense layers.
- **Attention Layers (Transformers):**
- The Self-Attention mechanism relies on the Softmax function.
- Softmax is a composition of exponentials and (positive) division, which is strictly real analytic everywhere.
- Therefore, interleaving Softmax attention layers with ReLU-based MLPs (as in modern Transformers) maintains the piecewise analytic property of the total loss landscape.

Conclusion

- The GD map is mathematically well-behaved even when the loss landscape is fractured by ReLUs.
- Pathological behaviors (collapsing entire regions of parameter space into saddle points) are mathematically impossible for almost all learning rates.

Thank You

Questions & Discussion